

Hello Soft Clustering (GMM)

T1. Using 3 mixtures, initialize your Gaussian with means (3,3), (2,2), and (-3,-3), and standard Covariance, I, the identity matrix. Use equal mixture weights as the initial weights. Repeat three iterations of EM. Write down $w_{n,j}$, m_j , $\vec{\mu}_j$, Σ_j for each EM iteration. (You may do the calculations by hand or write code to do so)

$$w_{n,j} = \frac{p(x_n; \vec{\mu}_j, \Sigma_j) m_j}{\sum_j p(x_n; \vec{\mu}_j, \Sigma_j) m_j} \quad (1)$$

$w_{n,j}$ means the probability that data point n comes from Gaussian number j .

Maximization: Update the model parameters, ϕ , $\vec{\mu}_j$, Σ_j .

$$m_j = \frac{1}{N} \sum_n w_{n,j} \quad (2)$$

$$\vec{\mu}_j = \frac{\sum_n w_{n,j} \vec{x}_n}{\sum_n w_{n,j}} \quad (3)$$

$$\Sigma_j = \frac{\sum_n w_{n,j} (\vec{x}_n - \vec{\mu}_j)(\vec{x}_n - \vec{\mu}_j)^T}{\sum_n w_{n,j}} \quad (4)$$

The above equation is used for full covariance matrices. For our small toy example, we will use diagonal covariance matrices, which can be acquired by setting the off-diagonal values to zero. In other words, $\Sigma_{(i,j)} = 0$, for $i \neq j$.

TODO: Complete functions below including

- Fill relevant parameters in each function.
- Implement computation and return values.

These functions will be used in T1-4.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt

# Hint: You can use this function to get gaussian distribution.
# https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.multiv
from scipy.stats import multivariate_normal
```

```
In [ ]: class GMM:
    def __init__(self, mixture_weight, mean_params, cov_params):
        """
        Initialize GMM.
        """
        # Copy construction values.
        self.mixture_weight = mixture_weight
```

```

self.mean_params = mean_params
self.cov_params = cov_params

# Initiaailize iteration.
self.n_iter = 0

def estimation_step(self, data, mixture_weight, mean_params, cov_param
prob = np.array([multivariate_normal(mean=mean_params[i], cov=cov
prob_m = np.dot(np.eye(mixture_weight.shape[0]) * mixture_weight,
w = prob_m / (np.sum(prob_m, axis=0))
return w

def maximization_step(self, data, w):
    """
    TODO: Perform maximization step.
    (Update parameters in this GMM model.)
    """
    self.mixture_weight = (1/data.shape[0])*np.sum(w, axis=1)

    self.mean_params = (np.dot(w, data).T / (w.sum(axis=1)+1e-10)).T

    self.cov_params = np.array([np.dot((w[i].reshape(-1,1))*(data-self
self.cov_params = (self.cov_params+1e-10) * np.eye(self.cov_param

def get_log_likelihood(self, data):
    """
    TODO: Compute log likelihood.
    """
    log_likelihood = np.log(np.array([multivariate_normal(mean=self.m

    return log_likelihood

def print_iteration(self):
    print("m :\n", self.mixture_weight)
    print("mu :\n", self.mean_params)
    print("covariance matrix :\n", self.cov_params)
    print("-----")

def perform_em_iterations(self, data, num_iterations, display=True):
    """
    Perform estimation & maximization steps with num_iterations.
    Then, return list of log_likelihood from those iterations.
    """
    log_prob_list = []

    # Display initialization.
    if display:
        print("Initialization")
        self.print_iteration()

    for n_iter in range(num_iterations):

        w = self.estimation_step(data, self.mixture_weight, self.mean
        self.maximization_step(data, w)

        # Calculate log prob.
        log_prob = self.get_log_likelihood(data)
        log_prob_list.append(log_prob)

```

```

        # Display each iteration.
        if display:
            print(f"Iteration: {n_iter}")
            self.print_iteration()

    return log_prob_list

```

```

In [ ]: # TODO
def plot_log_prob(log_prob_list):
    plt.figure(figsize=(20, 10))
    plt.grid()
    plt.title('Log likelihood history')
    plt.plot(np.arange(len(log_prob_list)), log_prob_list)
    plt.show()

```

T2. Plot the log likelihood of the model given the data after each EM step. In other words, plot $\log \prod_n p(\vec{x}_n | \varphi, \vec{\mu}, \Sigma)$. Does it go up every iteration just as we learned in class?

```

In [ ]: num_iterations = 3
num_mixture = 3
mixture_weight = np.array([1] * num_mixture) # m
mean_params = np.array([[3,3], [2,2], [-3,-3]], dtype = float)
cov_params = np.array([np.eye(2)] * num_mixture)

X, Y = np.array([1, 3, 2, 8, 6, 7, -3, -2, -7]), np.array([2, 3, 2, 8, 6,
data = np.vstack([X,Y]).T

gmm = GMM(mixture_weight, mean_params, cov_params)
log_prob_list_3 = gmm.perform_em_iterations(data, num_iterations)
plot_log_prob(log_prob_list_3)

```

Initialization

m :

[1 1 1]

mu :

[[3. 3.]

[2. 2.]

[-3. -3.]]

covariance matrix :

[[[1. 0.]

[0. 1.]]

[[1. 0.]

[0. 1.]]

[[1. 0.]

[0. 1.]]]

Iteration: 0

m :

[0.45757242 0.20909425 0.33333333]

mu :

[[5.78992692 5.81887265]

[1.67718211 2.14523106]

[-4. -4.66666666]]

covariance matrix :

[[[4.53619412 0.]

[0. 4.28700611]]

[[0.51645579 0.]

[0. 0.13152618]]

[[4.66666668 0.]

[0. 2.88888891]]]

Iteration: 1

m :

[0.40711618 0.25954961 0.33333421]

mu :

[[6.27176215 6.27262711]

[1.72091544 2.14764812]

[-3.99998589 -4.6666488]]

covariance matrix :

[[[2.94672736 0.]

[0. 2.93847196]]

[[0.49649261 0.]

[0. 0.12584815]]

[[4.66673088 0.]

[0. 2.88900236]]]

Iteration: 2

m :

[0.36070909 0.30595677 0.33333414]

mu :

[[6.6962644 6.69629468]

[1.91071238 2.27383436]

[-3.99998673 -4.6666501]]

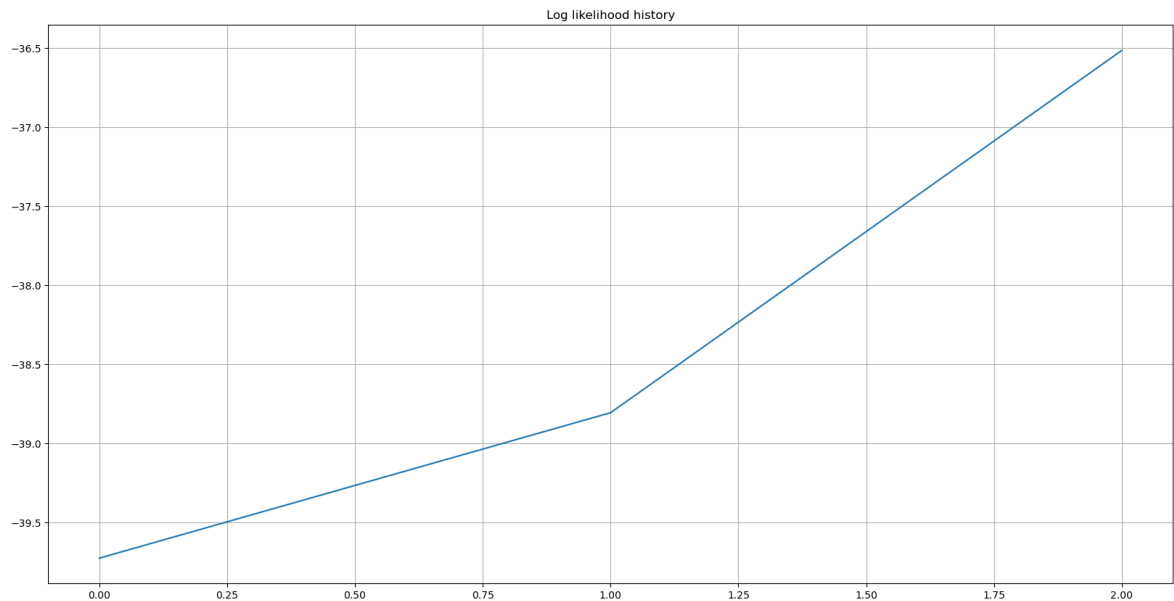
covariance matrix :

[[[1.73961067 0.]

```
[0.          1.73929602]]
```

```
[[0.62898406 0.          ]  
 [0.          0.1988491  ]]
```

```
[[4.66672942 0.          ]  
 [0.          2.88899545]]]
```



ANS : Yes it looklike going to coverge every iteration

T3. Using 2 mixtures, initialize your Gaussian with means (3,3) and (-3,-3), and standard Covariance, I, the identity matrix. Use equal mixture weights as the initial weights. Repeat three iterations of EM. Write down $w_{n,j}$, m_j , $\vec{\mu}_j$, Σ_j for each EM iteration.

```
In [ ]: num_mixture = 2  
mixture_weight = [1] * num_mixture  
mixture_weight = np.array(mixture_weight)  
  
mean_params = np.array([[3,3], [-3,-3]], dtype = float)  
cov_params = np.array([np.eye(2)] * num_mixture)  
  
# INSERT CODE HERE  
gmm = GMM(mixture_weight, mean_params, cov_params)  
log_prob_list_2 = gmm.perform_em_iterations(data, num_iterations)  
  
plot_log_prob(log_prob_list_2)
```

Initialization

m :

[1 1]

mu :

[[3. 3.]

[-3. -3.]]

covariance matrix :

[[[1. 0.]

[0. 1.]]

[[1. 0.]

[0. 1.]]]

Iteration: 0

m :

[0.66666666 0.33333334]

mu :

[[4.50000001 4.66666667]

[-3.99999997 -4.66666663]]

covariance matrix :

[[[6.91666665 0.]

[0. 5.88888889]]

[[4.66666677 0.]

[0. 2.8888891]]]

Iteration: 1

m :

[0.66669436 0.33330564]

mu :

[[4.49961311 4.66620178]

[-3.99993241 -4.66651231]]

covariance matrix :

[[[6.91944755 0.]

[0. 5.89275124]]

[[4.66806942 0.]

[0. 2.89103318]]]

Iteration: 2

m :

[0.66669453 0.33330547]

mu :

[[4.49961084 4.66619903]

[-3.99993206 -4.6665114]]


covariance matrix :

[[[6.91946372 0.]

[0. 5.8927741]]

[[4.66807754 0.]

[0. 2.89104566]]]




T4. Plot the log likelihood of the model given the data after each EM step. Compare the log likelihood between using two mixtures and three mixtures. Which one has the better likelihood?

```
In [ ]: # TODO: Plot Comparision of log_likelihood from T1 and T3

plt.figure(figsize=(20, 10))
plt.plot(np.arange(len(log_prob_list_2)), log_prob_list_2, label='2 Mixtu
plt.scatter(np.arange(len(log_prob_list_2)), log_prob_list_2, marker='x',
plt.plot(np.arange(len(log_prob_list_3)), log_prob_list_3, label='3 Mixtu
plt.scatter(np.arange(len(log_prob_list_3)), log_prob_list_3, marker='o',

plt.legend()
plt.grid()
```

[-39.725993156389336, -38.80573197942934, -36.512976308824065]



ANS : Likelihood with 3 mixture is better Since it converge.

The face database

```
In [ ]: # Download facedata for google colab
# !wget -nc https://github.com/ekapolc/Pattern_2024/raw/main/HW/HW03/face
# !unzip facedata_mat.zip
```

```
In [ ]: import scipy.io
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from skimage import img_as_float

# Change path to your facedata.mat file.
facedata_path = 'facedata.mat'

data = scipy.io.loadmat(facedata_path)
data_size = data['facedata'].shape

%matplotlib inline
data_size
```

```
Out[ ]: (40, 10)
```

Preprocess xf

```
In [ ]: xf = np.zeros((data_size[0], data_size[1], data['facedata'][0,0].shape[0])
for i in range(data['facedata'].shape[0]):
    for j in range(data['facedata'].shape[1]):
        xf[i,j] = img_as_float(data['facedata'][i,j])
```

```
In [ ]: # Example: Plotting face image.
plt.figure(figsize=(20, 10))
for i in range(5):
    for j in range(10):
        plt.subplot(5, 10, i*10+j+1)
        plt.imshow(xf[i,j], cmap='gray')
        plt.title(f'(Face {i},{j})')
        plt.axis('off')
```




T5. What is the Euclidean distance between $xf[0,0]$ and $xf[0,1]$? What is the Euclidean distance between $xf[0,0]$ and $xf[1,0]$? Does the numbers make sense? Do you think these numbers will be useful for face verification?

```
In [ ]: def L2_dist(x1, x2):
         distance = np.sqrt(np.sum((x1 - x2)**2))
         return distance

# Test L2_dist
def test_L2_dist():
    assert L2_dist(np.array([1, 2, 3]), np.array([1, 2, 3])) == 0.0
    assert L2_dist(np.array([0, 0, 0]), np.array([1, 2, 3])) == np.sqrt(10)

test_L2_dist()


print('Euclidean distance between xf[0,0] and xf[0,1] is', L2_dist(xf[0,0], xf[0,1]))
print('Euclidean distance between xf[0,0] and xf[1,0] is', L2_dist(xf[0,0], xf[1,0]))
```

Euclidean distance between $xf[0,0]$ and $xf[0,1]$ is 10.037616294165492
 Euclidean distance between $xf[0,0]$ and $xf[1,0]$ is 8.173295099737281

```
In [ ]: # TODO: Show why does the numbers make sense
print('Euclidean distance between xf[0,0] and xf[0,0] is', L2_dist(xf[0,0], xf[0,0]))
print('Euclidean distance between xf[0,0] and xf[0,1] is', L2_dist(xf[0,0], xf[0,1]))
print('Euclidean distance between xf[1,0] and xf[0,1] is', L2_dist(xf[1,0], xf[0,1]))
print('Euclidean distance between xf[1,0] and xf[0,0] is', L2_dist(xf[1,0], xf[0,0]))

plt.subplot(1,3,1)
plt.imshow(xf[0,0], cmap = 'gray')
plt.subplot(1,3,2)
plt.imshow(xf[0,1], cmap = 'gray')
plt.subplot(1,3,3)
plt.imshow(xf[1,0], cmap = 'gray')
plt.show()
```

Euclidean distance between $xf[0,0]$ and $xf[0,0]$ is 0.0
 Euclidean distance between $xf[0,0]$ and $xf[0,1]$ is 10.037616294165492
 Euclidean distance between $xf[1,0]$ and $xf[0,1]$ is 11.517134319336666
 Euclidean distance between $xf[1,0]$ and $xf[0,0]$ is 8.173295099737281



ANS : Euclidian values seems make sense when comparing difference of images in overall. This can apply only when image has same dimension. It's sensitive to rotation and distance of objects in images so it might not good for facial recognition.

T6. Write a function that takes in a set of feature vectors T and a set of feature vectors D, and then output the similarity matrix A. Show the matrix as an image. Use the feature vectors from the first 3 images from all 40 people for list T (in order $x[0, 0]$, $x[0, 1]$, $x[0, 2]$, $x[1, 0]$, $x[1, 1]$, ... $x[39, 2]$). Use the feature vectors from the remaining 7 images from all 40 people for list D (in order $x[0, 3]$, $x[0, 4]$, $x[0, 5]$, $x[1, 6]$, $x[0, 7]$, $x[0, 8]$, $x[0, 9]$, $x[1, 3]$, $x[1, 4]$... $x[39, 9]$). We will treat T as our training images and D as our testing images

```
In [ ]: def organize_shape(matrix):
    """
    TODO (Optional): Reduce matrix dimension of 2D image to 1D and merge
    This function can be useful at organizing matrix shapes.

    Example:
        Input shape: (people_index, image_index, image_shape[0], image_shape[1])
        Output shape: (people_index*image_index, image_shape[0]*image_shape[1])
    """

    people_index, image_index, image_shape_0, image_shape_1 = matrix.shape

    # Reshape the matrix to merge people and image dimensions
    reshaped_matrix = matrix.copy().reshape(people_index * image_index, image_shape_0 * image_shape_1)

    return reshaped_matrix

def generate_similarity_matrix(A, B):
    """
    TODO: Calculate similarity matrix M,
    which  $M[i, j]$  is a distance between  $A[i]$  and  $B[j]$ .
    """
```

```

distance_matrix = np.sqrt(np.sum((A[:, np.newaxis] - B)**2, axis=2))

return distance_matrix

def test_generate_similarity_matrix():
    test_A = np.array([[1, 2], [3, 4]])
    test_B = np.array([[1, 2], [5, 6], [7, 8]])
    expected_matrix = np.sqrt(np.array([[0, 32, 72], [8, 8, 32]]))
    assert (generate_similarity_matrix(test_A, test_B) == expected_matrix)

test_generate_similarity_matrix()

```

```

In [ ]: #TODO: Show similarity matrix between T and D.
T = organize_shape(xf[:, :3])
D = organize_shape(xf[:, 3:])
similarity_matrix = generate_similarity_matrix(organize_shape(xf[:, :3]),
similarity_matrix

```

```

Out[ ]: array([[10.36960631,  9.84869463,  8.99622801, ...,  9.89638826,
           9.36151948, 10.66062617],
        [11.24987522,  7.4172114 ,  9.88066979, ..., 10.9694266 ,
          10.90238961, 10.93630575],
        [10.22209276,  9.41321639,  9.29987742, ...,  9.99689456,
          9.94499521, 10.41008147],
        ...,
        [11.50315003, 10.77719551, 11.45361976, ...,  5.5842077 ,
          9.25718815,  7.93519019],
        [ 9.82401314, 10.09262219,  9.9006693 , ...,  6.75074775,
          7.45046954,  8.90196597],
        [ 9.75449186,  9.98745234, 10.16210857, ...,  6.51894822,
          7.88934623,  9.55831251]])

```

T7. From the example similarity matrix above, what does the black square between [5:10,5:10] suggest about the pictures from person number 2? What do the patterns from person number 1 say about the images from person 1?

```

In [ ]: # INSERT CODE HERE
# Show the similarity matrix
plt.imshow(similarity_matrix, cmap='gray')
plt.colorbar()
plt.title("Similarity Matrix")
plt.show()

# Extract the relevant squares
person_2_square = similarity_matrix[5:10, 5:10]
person_1_square = similarity_matrix[:3, :3]

# Calculate the average similarity for each square
person_2_avg_similarity = np.mean(person_2_square)
person_1_avg_similarity = np.mean(person_1_square)

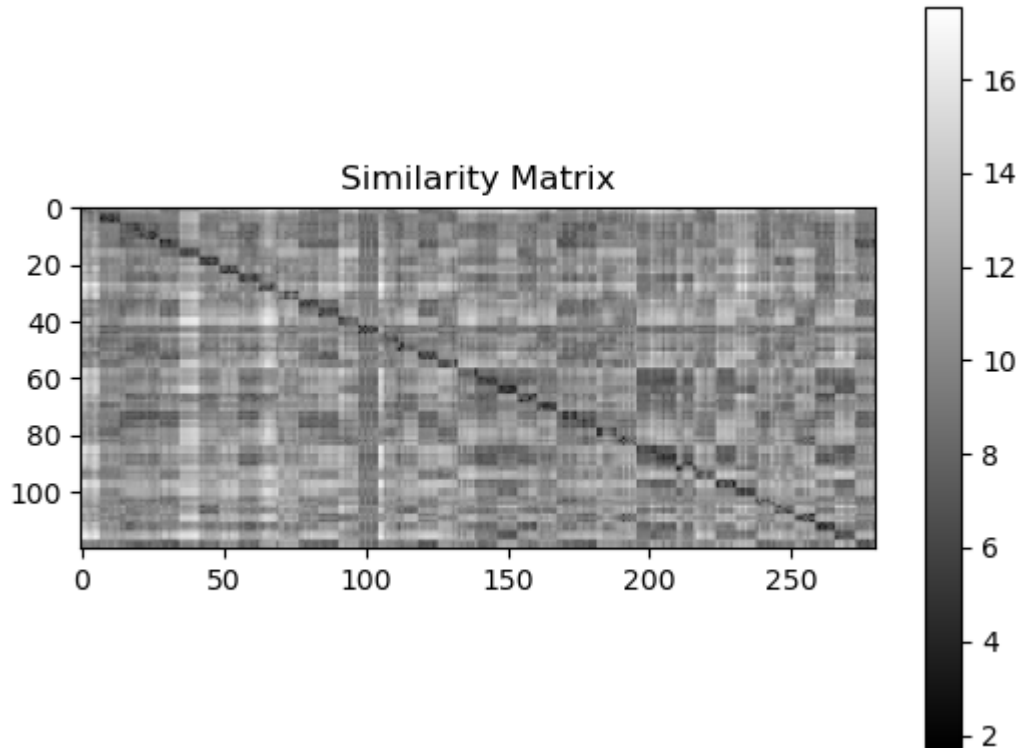
print("Average similarity for person number 2:", person_2_avg_similarity)
print("Average similarity for person number 1:", person_1_avg_similarity)

eT = organize_shape(xf[:5, :5])


```

```
plt.imshow(generate_similarity_matrix(eT, eT), cmap='gray')

plt.figure(figsize=(20, 10))
for i in range(2):
    for j in range(10):
        plt.subplot(2, 10, i*10+j+1)
        plt.imshow(xf[i,j], cmap='gray')
        plt.title(f'Face {i},{j}')
        plt.axis('off')
```



Average similarity for person number 2: 9.637038405598823
 Average similarity for person number 1: 9.633052437208354



ANS : The black square between [5:10,5:10] in the similarity matrix suggests that the images from person number 2 have a low similarity with each other. This means that the images of person number 2 have significant variations or differences among them. On the other hand, the patterns from person number 1 in the similarity matrix suggest that the images from person number 1 have a high similarity with each other. This indicates that the images of person number 1 have a consistent pattern or similarity among them.

T8. Write a function that takes in the similarity matrix created from the previous part, and a threshold t as inputs. The outputs of the function are the true positive rate and the false alarm rate of the face verification task (280 Test images, tested on 40 people, a total of 11200

testing per threshold). What is the true positive rate and the false alarm rate for $t = 10$?

```
In [ ]: def evaluate_performance(similarity_matrix, threshold):
        """
        TODO: Calculate true positive rate and false alarm rate from given si
        """
        y_pred = np.zeros((40, 280))
        y_actual = np.zeros((40, 280))
        for i in range(40):
            for j in range(280):
                y_pred[i, j] = similarity_matrix[3*i:3*i+3, j].min() < thresh
                y_actual[i, j] = i == j//7

        tp = np.where(y_pred==y_actual, y_pred, 0).sum()
        tn = np.where(y_pred==y_actual, 1-y_pred, 0).sum()
        fp = np.where(y_pred!=y_actual, 1-y_actual, 0).sum()
        fn = np.where(y_pred!=y_actual, y_actual, 0).sum()

        true_pos_rate = tp/(tp+fn)
        false_alarm_rate = fp/(tn+fp)

        return true_pos_rate, false_alarm_rate

# Quick check
# (true_pos_rate, false_neg_rate) should be (0.9928571428571429, 0.335073
evaluate_performance(similarity_matrix, 9.5)
```

```
Out[ ]: (0.9928571428571429, 0.33507326007326005)
```

```
In [ ]: # INSERT CODE HERE
evaluate_performance(similarity_matrix, 10)
```

```
Out[ ]: (0.9964285714285714, 0.4564102564102564)
```

ANS: True Positive rate is 0.9964285714285714 and False Alarm is 0.4564102564102564

T9. Plot the RoC curve for this simple verification system. What should be the minimum threshold to generate the RoC curve? What should be the maximum threshold? Your RoC should be generated from at least 1000 threshold levels equally spaced between the minimum and the maximum. (You should write a function for this).

```
In [ ]: def calculate_roc(input_mat):
        """
        TODO: Calculate a list of true_pos_rate and a list of false_neg_rate
        """
        tpr_list = []
        far_list = []

        n_iter = 2000
```

```

recall = None
diff_recal = float('inf')
eer = 0
diff = float('inf')

rag = input_mat.max() - input_mat.min()
minv = input_mat.min() + 0.20*rag
maxv = input_mat.max() - 0.20*rag
for threshold in np.linspace(minv, maxv, n_iter):
    tpr, far = evaluate_performance(input_mat, threshold)
    if abs(far-0.001) < diff_recal:
        recall = tpr
        diff_recal = abs(far-0.001)
    if abs(tpr+far-1) < diff:
        diff = abs(tpr+far-1)
        eer = tpr
    tpr_list.append(tpr)
    far_list.append(far)

return tpr_list, far_list, eer, recall

def plot_roc(input_mat, label=None):
    """
    TODO: Plot RoC Curve from a given matrix.
    """
    tpr_list, far_list, eer, recall = calculate_roc(input_mat)
    plt.title("Plot of ROC")
    if label:
        plt.plot(far_list, tpr_list, label=label)
    else:
        plt.plot(far_list, tpr_list)
    plt.xlabel('False Alarm Rate')
    plt.ylabel('True Positive Rate')
    print(f"EER : {eer}")
    print(f"Recall at 0.1% FAR : {recall}")

    return eer

```

```

In [ ]: # INSERT CODE HERE
        plot_roc(similarity_matrix)

```

```


EER : 0.9107142857142857
Recall at 0.1% FAR : 0.5464285714285714

```

```

Out[ ]: 0.9107142857142857

```



ANS: I decide to choose minimum and maximum threshold by $\text{min/max} \pm 0.20 * \text{range}$

T10. What is the EER (Equal Error Rate)? What is the recall rate at 0.1% false alarm rate? (Write this in the same function as the previous question)

```
In [ ]: # You can add more parameter(s) to the function in the previous question.

# EER should be either 0.9071428571428571 or 0.9103759398496248 depending
# Recall rate at 0.1% false alarm rate should be 0.5428571428571428.
def plot_roc(input_mat, display=True, with_err=True, label=None):
    """
    Plot RoC Curve from a given matrix.
    """
    tpr_list, far_list, err, recall = calculate_roc(input_mat)

    # Calculate recall rate at FAR

    if display and with_err:
        plt.plot(1-np.array(tpr_list), tpr_list, color='gray', linestyle='dashed')

    # Plot ROC curve
    if display:
        plt.plot(far_list, tpr_list, label=label if label else 'ROC Curve')
        plt.xlabel('FAR')
        plt.ylabel('TPR')
        plt.title('ROC Curve')
        plt.legend()
```




```

plt.show()
print("ERR:", err)
print("Recall at FAR:", recall)
return err, recall

# Plot ROC curve with EER and recall rate at 0.1% FAR
plot_roc(similarity_matrix)

```



```

ERR: 0.9107142857142857
Recall at FAR: 0.5464285714285714

```

```

Out[ ]: (0.9107142857142857, 0.5464285714285714)

```

ANS:

T11. Compute the mean vector from the training images. Show the vector as an image (use `numpy.reshape()`). This is typically called the meanface (or meanvoice for speech signals). Your answer should look exactly like the image shown below.

```

In [ ]: # INSERT CODE HERE
meanface = T.mean(axis=0)

plt.title('mean face of training face')
plt.axis('off')
plt.imshow(meanface.reshape(56, 46), cmap='gray')
plt.show()

```

mean face of training face



T12. What is the size of the covariance matrix? What is the rank of the covariance matrix?

```
In [ ]: # TODO: Find the size and the rank of the covariance matrix.
print(f"""
Covariance matrix size : ({T.shape[1]},{T.shape[1]})
Rank of Covariance matrix : {min(T.shape[1], T.shape[0]-1)}
""").strip()
```

```
Covariance matrix size : (2576,2576)
Rank of Covariance matrix : 119
```

ANS:

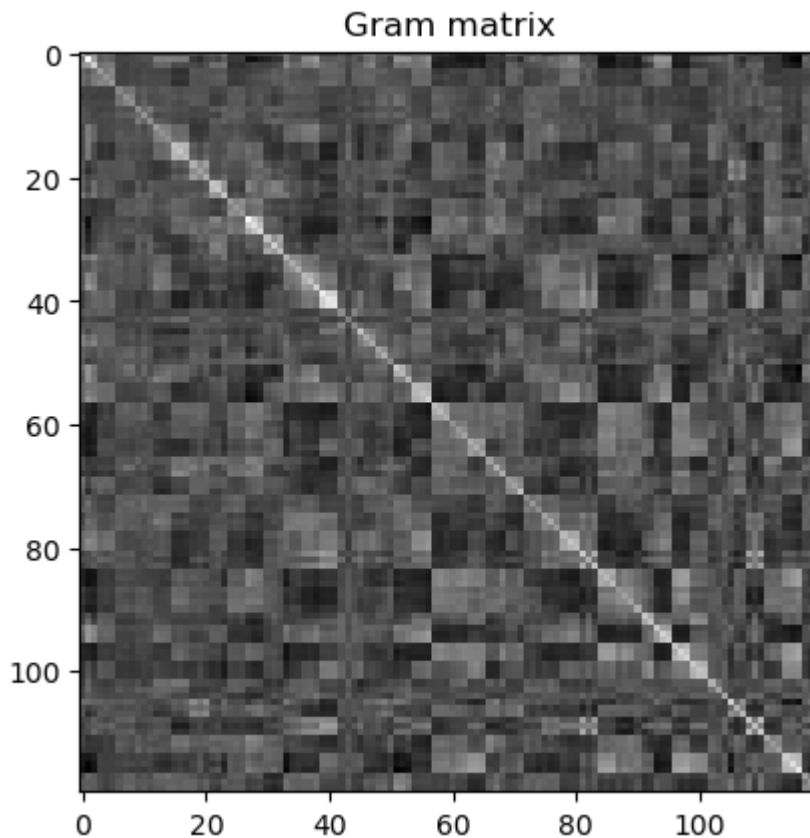
T13. What is the size of the Gram matrix? What is the rank of Gram matrix? If we compute the eigenvalues from the Gram matrix, how many non- zero eigenvalues do we expect to get?

```
In [ ]: # TODO: Compute gram matrix.
gram_matrix = np.matmul(T-meanface, (T-meanface).T)

plt.title(f'Gram matrix')
plt.imshow(gram_matrix, cmap='gray')

print(f"""
Size of gram matrix : {gram_matrix.shape}
Rank of gram matrix : {gram_matrix.shape[0]}
Expected non-zero eigenvalues : {gram_matrix.shape[0]-1}
""").strip()
```

Size of gram matrix : (120, 120)
Rank of gram matrix : 120
Expected non-zero eigenvalues : 119



ANS:

T14. Is the Gram matrix also symmetric? Why?

ANS:

T15. Compute the eigenvectors and eigenvalues of the Gram matrix, v_0 and λ . Sort the eigenvalues and eigenvectors in descending order so that the first eigenvalue is the highest, and the first eigenvector corresponds to the best direction. How many non-zero eigenvalues are there? If you see a very small value, it is just numerical error and should be treated as zero.

```
In [ ]: # Hint: https://numpy.org/doc/stable/reference/generated/numpy.linalg.eig

def calculate_eigenvectors_and_eigenvalues(matrix):
    """
    TODO: Calculate eigenvectors and eigenvalues,
    then sort the eigenvalues and eigenvectors in descending order.

    Hint: https://numpy.org/doc/stable/reference/generated/numpy.linalg.e
    """
```

```

# INSERT CODE HERE
eigenvalues, eigenvectors = np.linalg.eigh(matrix)

# Sort the eigenvalues and eigenvectors in descending order
idx = np.argsort(eigenvalues)[::-1]
eigenvalues = eigenvalues[idx]
eigenvectors = eigenvectors[:, idx]

# Treat very small values as zero
eigenvalues[np.abs(eigenvalues) < 1e-20] = 0.0
eigenvectors[np.abs(eigenvectors) < 1e-20] = 0.0

return eigenvalues, eigenvectors

eigenvalues, eigenvectors = calculate_eigenvectors_and_eigenvalues(gram_m

def test_eigenvalues_eigenvectors():
    # Dot product of an eigenvector pair should equal to zero.
    assert np.round(eigenvectors[10].dot(eigenvectors[20]), 10) == 0.0

    # Check if eigenvalues are sorted.
    assert list(eigenvalues) == sorted(eigenvalues, reverse = True)

test_eigenvalues_eigenvectors()

```

```

In [ ]: gram_evalues, gram_evectors = calculate_eigenvectors_and_eigenvalues(gram
print(f"""
Amount of non-zero eigenvalues of gram matrix : {np.sum(gram_evalues>1e-2
""").strip())

```

Amount of non-zero eigenvalues of gram matrix : 119

ANS:

T16. Plot the eigenvalues. Observe how fast the eigenvalues decrease. In class, we learned that the eigenvalues is the size of the variance for each eigenvector direction. If I want to keep 95% of the variance in the data, how many eigenvectors should I use?

```

In [ ]: # INSERT CODE HERE
total_variance = gram_evalues.sum()

top95_idx = 0
curr_var = 0
while curr_var < total_variance*0.95 or top95_idx >= len(gram_evalues):
    curr_var += gram_evalues[top95_idx]
    top95_idx += 1
top95_idx -= 1


print(f"""
Total variance : {total_variance:3,.2f}
95% of variance use {top95_idx+1} eigenvalues
""").strip()

```

Total variance : 6,853.80
95% of variance use 64 eigenvalues

```
In [ ]: plt.title("Plot of eigenvalues from gram matrix")
plt.plot(gram_evalues)
plt.vlines(top95_idx,
           ymin=0,
           ymax=gram_evalues[top95_idx],
           color='r',
           linestyle='--',
           label='95% of variance')
plt.fill_between(np.arange(top95_idx+1),
                 gram_evalues[:top95_idx+1],
                 color='r')
plt.legend()
```

Out[]: <matplotlib.legend.Legend at 0x7950e7285f70>



ANS:

T17. Compute \vec{v} . Don't forget to renormalize so that the norm of each vector is 1 (you can use `numpy.linalg.norm`). Show the first 10 eigenvectors as images. Two example eigenvectors are shown below. We call these images eigenfaces (or eigenvoice for speech signals).


```
In [ ]: # TODO: Compute v, then renormalize it.
v = np.matmul((T-meanface).T, gram_evecs)
v = v/np.linalg.norm(v, axis=0)
```

```
v.shape
```

```
Out[ ]: (2576, 120)
```

```
In [ ]: def test_eigenvector_cov_norm(v):  
        assert (np.round(np.linalg.norm(v, axis=0), 1) == 1.0).all()  
  
        test_eigenvector_cov_norm(v)
```

```
In [ ]: # TODO: Show the first 10 eigenvectors as images.  
plt.figure(figsize=(20, 10))  
for i in range(10):  
    plt.subplot(2, 5, i+1)  
    plt.title(f"eigenface {i}")  
    plt.imshow(v[:, i].reshape(56,46), cmap='gray')
```



T18. From the image, what do you think the first eigenvector captures? What about the second eigenvector? Look at the original images, do you think biggest variance are capture in these two eigenvectors?

ANS:

T19. Find the projection values of all images. Keep the first $k = 10$ projection values. Repeat the simple face verification system we did earlier using these projected values. What is the EER and the recall rate at 0.1% FAR?

```
In [ ]: def calculate_projection_vectors(matrix, meanface, v, k):  
        """  
        TODO: Find the projection vectors on v from given matrix and meanface  
        """  
        projection_vectors = np.matmul(matrix-meanface, v[:, :k])  
  
        return projection_vectors
```

```

In [ ]: # TODO: Get projection vectors of T and D, then Keep first k projection v
k = 10
T_reduced = calculate_projection_vectors(T, meanface, v, k)
D_reduced = calculate_projection_vectors(D, meanface, v, k)


def test_reduce_dimension():
    assert T_reduced.shape[-1] == k
    assert D_reduced.shape[-1] == k

test_reduce_dimension()

In [ ]: # TODO: Get similarity matrix of T_reduced and D_reduced
reduced_similarity_matrix = generate_similarity_matrix(T_reduced, D_reduced)

In [ ]: # TODO: Find EER and the recall rate at 0.1% FAR.
plot_roc(reduced_similarity_matrix)

```



ERR: 0.9214285714285714
Recall at FAR: 0.75

ANS:

T20. What is the k that gives the best EER? Try k = 5, 6, 7, 8, 9, 10, 11, 12, 13, 14.

```

In [ ]: # INSERT CODE HERE
best_eer = -1
best_k = None
for k in range(5, 15):

```

```


print(f"for k = {k}")
T_reduced = calculate_projection_vectors(T, meanface, v, k)
D_reduced = calculate_projection_vectors(D, meanface, v, k)

reduced_similarity_matrix = generate_similarity_matrix(T_reduced, D_r
err, recall =plot_roc(reduced_similarity_matrix)
if best_eer == -1 or best_eer > err:
    best_eer = err
    best_k = k

plt.legend()
print(f"Best k : {best_k}")

```


for k = 5



ERR: 0.8928571428571429

Recall at FAR: 0.8428571428571429


for k = 6



ERR: 0.9071428571428571

Recall at FAR: 0.8


for $k = 7$




ERR: 0.9071428571428571

Recall at FAR: 0.7857142857142857


for $k = 8$



ERR: 0.9142857142857143
 Recall at FAR: 0.775
 for $k = 9$




ERR: 0.9178571428571428
 Recall at FAR: 0.7678571428571429
 for $k = 10$



ERR: 0.9214285714285714

Recall at FAR: 0.75


for $k = 11$



ERR: 0.9214285714285714

Recall at FAR: 0.7214285714285714


for $k = 12$



ERR: 0.9142857142857143

Recall at FAR: 0.7


for $k = 13$



ERR: 0.9178571428571428

Recall at FAR: 0.6857142857142857

for $k = 14$




No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

ERR: 0.9178571428571428

Recall at FAR: 0.6642857142857143

Best k : 5



ANS:

T21. In order to assure that S_W is invertible we need to make sure that S_W is full rank. How many PCA dimensions do we need to keep in order for S_W to be full rank? (Hint: How many dimensions does S_W have? In order to be of full rank, you need to have the same number of linearly independent factors)

ANS:

```
In [ ]: # TODO: Define dimension of PCA.
n_dim = T.shape[0] - 40
print(f"Dimension : {n_dim}")

# TODO: Find PCA of T and D with n_dim dimension.
T_reduced = calculate_projection_vectors(T, meanface, v, n_dim)
D_reduced = calculate_projection_vectors(D, meanface, v, n_dim)

print(T_reduced.shape, D_reduced.shape)

# TODO: Find PCA of T and D with n_dim dimension.
```

```
Dimension : 80
(120, 80) (280, 80)
```

T22. Using the answer to the previous question, project the original input to the PCA subspace. Find the LDA projections. To find the inverse, use `-1 numpy.linalg.inv`. Is S_W S_B symmetric? Can we still use `numpy.linalg.eigh`? How many non-zero eigenvalues are there?

```
In [ ]: # TODO: Find the LDA projection.
T_reduced_by_class = T_reduced.reshape((40, 3, -1))
class_mean = T_reduced_by_class.mean(axis=1)
all_mean = class_mean.mean(axis=0).reshape(1, -1)

s_b = np.array([np.matmul((class_mean[i]-all_mean).T, (class_mean[i]-all_

s_wi = np.array([np.matmul((T_reduced_by_class[i]-class_mean[i]).T, (T_re
s_w = s_wi.sum(axis=0)

LDA = np.matmul(np.linalg.inv(s_w), s_b)
print(f"LDA is symmetric? : {np.allclose(LDA, LDA.T)}")
LDA_evalues, LDA_evectors = np.linalg.eig(LDA)

LDA_evectors = LDA_evectors.real
LDA_evalues = LDA_evalues.real
```

```
LDA is symmetric? : False
```

ANS:

T23. Plot the first 10 LDA eigenvectors as images (the 10 best projections). Note that in this setup, you need to convert back to the original image space by using the PCA projection. The LDA eigenvectors can be considered as a linear combination of eigenfaces. Compare the LDA projections with the PCA projections.

```
In [ ]: # INSERT CODE HERE
print(f"Amount of non-zero eigenvalues : {np.where(LDA_evalues>1e-5, 1, 0
```

Amount of non-zero eigenvalues : 39

T24. The combined PCA+LDA projection procedure is called fisherface. Calculate the fisherfaces projection of all images. Do the simple face verification experiment using fisherfaces. What is the EER and recall rate at 0.1% FAR?


```
In [ ]: # INSERT CODE HERE
best_10_LDA = LDA_evalues[:, :10]
T_LDA = np.matmul(T_reduced, best_10_LDA)
T_eigenface = np.matmul(v[:, :n_dim], best_10_LDA)

D_LDA = np.matmul(D_reduced, best_10_LDA)
D_eigenface = np.matmul(v[:, :n_dim], best_10_LDA)

plt.figure(figsize=(20, 10))
for idx in range(10):
    plt.subplot(2, 5, idx+1)
    plt.imshow(v[:, idx].reshape(56, 46), cmap='gray')
    plt.title(f"LDA {idx}")
    plt.axis('off')

plt.show()

T_LDA.shape, D_LDA.shape
```




Out[]: ((120, 10), (280, 10))

ANS:

T25.Plot the RoC of all three experiments (No projection, PCA, andFisher) on the same axes. Compare and contrast the three results. Submit yourwriteup and code on MyCourseVille.

```
In [ ]: # INSERT CODE HERE
reduced_similarity_matrix = generate_similarity_matrix(T_LDA, D_LDA)
eer = plot_roc(reduced_similarity_matrix, label='LDA')
```

ERR: 0.8857142857142857
Recall at FAR: 0.5642857142857143

ANS: